

AICOSS - AI Special Program

Sample solutions to exercise for data cleaning

In the original [Kaggle Competition](#) for this data set, a prediction model was to be developed for the variable `price_doc` (*target*). In order to train such a model, however, some variables must first be adjusted. This is what you will do in this exercise.

Clean up the following variables:

- **state**
`state` indicates the status of the apartment and should contain the values 1-4.
- **floor and max_floor**
What is the connection between `floor` and `max_floor`? Find and correct any inconsistencies.
- **build_year**
Check `build_year` for incorrect values and correct them in a meaningful way.
- **life_sq**
Remove outlier from `life_sq` and perform two different types of imputation.

```
In [1]: import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df = pd.read_csv('./data/sberbank.csv')
```

state

```
In [3]: df['state'].unique()
```

```
Out[3]: array([nan,  3.,  1.,  2.,  4., 33.])
```

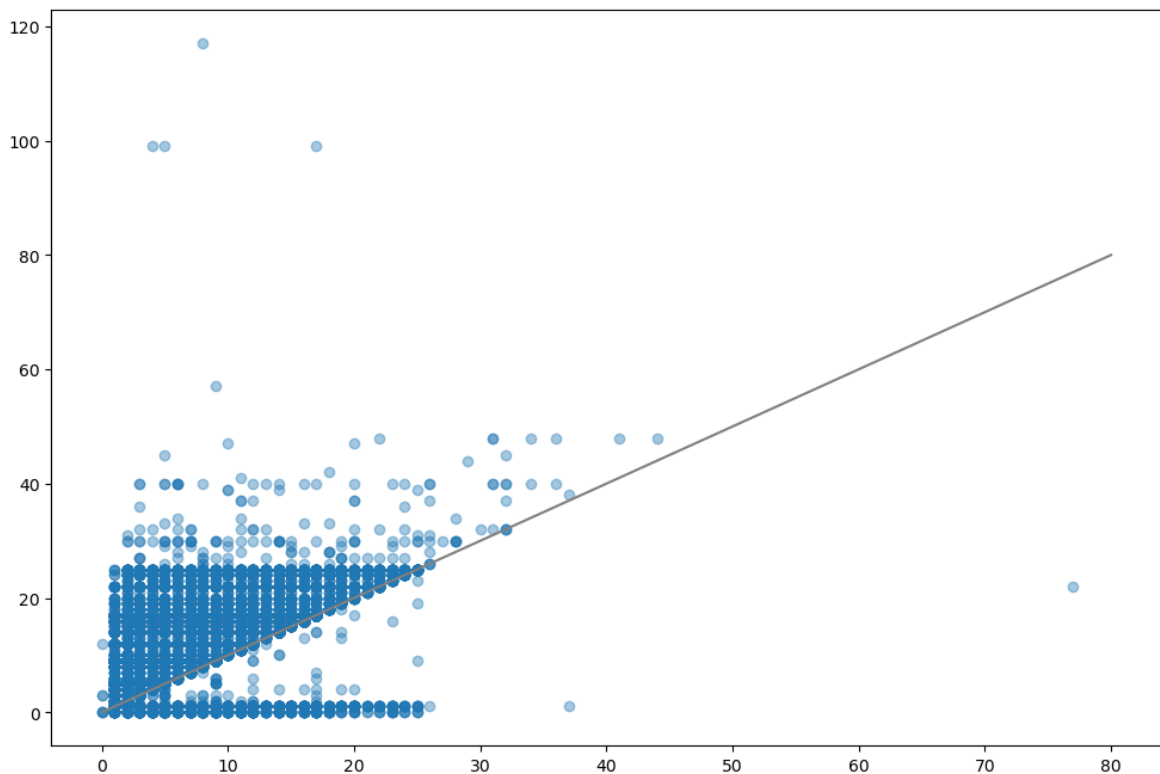
The 33rd could be a typing error. We replace with 3.:

```
In [4]: df['state'].replace(33., 3., inplace=True)
```

floor und max_floor

Check the value combinations:

```
In [5]: f, ax = plt.subplots(figsize=(12, 8))
plt.scatter(x=df['floor'], y=df['max_floor'], alpha=0.4)
plt.plot([0, 80], [0, 80], color='.5');
```



Actually, we should not find any values below the diagonal, as the number of maximum storeys is then smaller than the number of existing storeys.

```
In [6]: df.loc[df['max_floor'] < df['floor'],:].shape
```

```
Out[6]: (1493, 292)
```

This is the case for 1493 entries! We can replace the incorrect values (as well as the zero values) with `NaN` :

```
In [7]: mask = (df['max_floor'] < df['floor']) | (df['max_floor']==0.0)
```

```
In [8]: df.loc[mask, 'max_floor'] = np.NaN
```

build_year

The following values are available in `build_year` :

```
In [9]: df['build_year'].dropna().astype('int64').unique()
```

```
Out[9]: array([[ 1907,   1980,   2014,   1970,   1982,   2013,
    2004,   2003,   1957,   1986,   1960,   1995,
    1979,   1975,   1987,   1962,   1969,   1993,
    1996,   1972,   2011,   1965,   2010,   1985,
    2006,   1961,   1971,   1978,   1966,   1967,
    2000,   1964,   1977,   1983,   1968,   1974,
    2008,   1959,   2007,   1984,   1976,   1997,
    1989,   1958,   1988,   2012,   1990,   1946,
    1917,   2002,   2005,   2001,   1963,   1954,
    1951,   1981,   1955,   1999,   2009,   1973,
    1994,   1998,   1992,   1950,   1956,   2015,
     0,   1932,     1,   1937,   1938,   1939,
    1991,   1934,   1935, 20052009,   1947,   1953,
    1933,   2016,   1930,   1912,   1929,     3,
    1928,   1915,   1936,   1925,   1940,   1943,
    1927,   1896,   1911,   1924,   1952,   2017,
    1926,   1931,    20,   1860,   1949,   1914,
    4965,   1910,   1895,   1948,   1876,   1900,
    1890,   1920,   1904,   1906,   1941,   1691,
    1905,   1886,   2018,   215,     71])
```

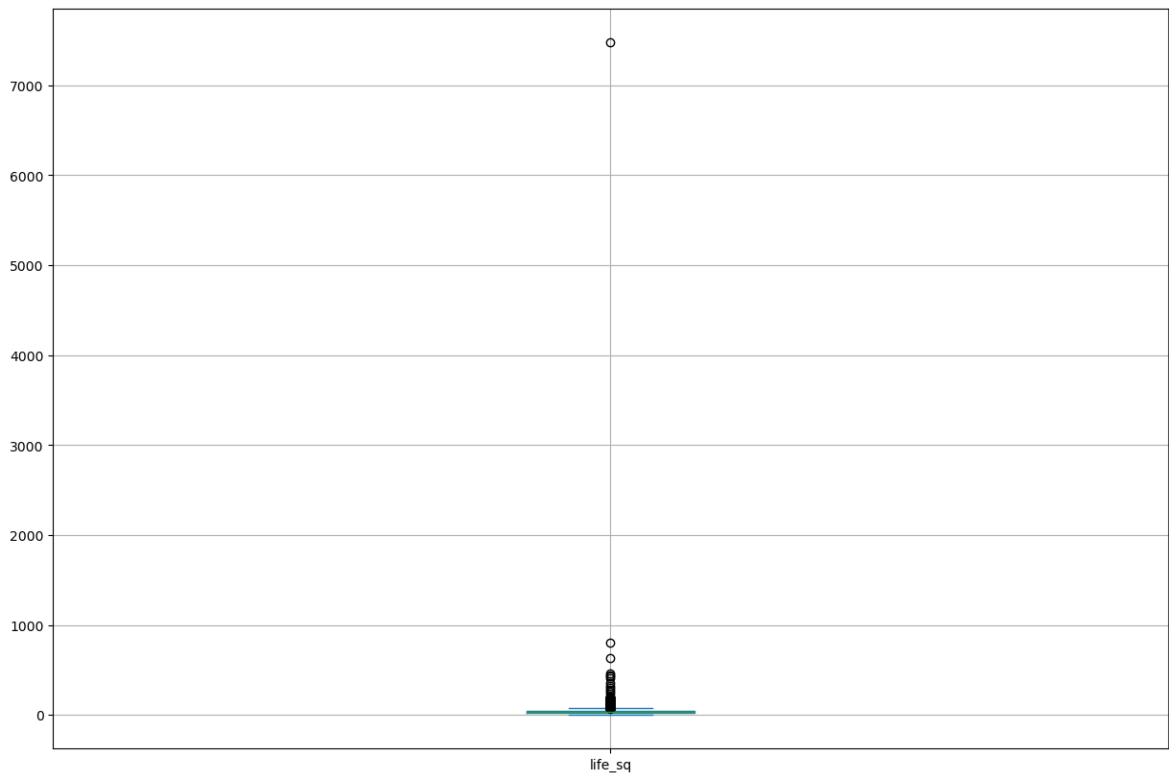
The following values must be replaced:

- 0 with `2000
- 1 with `2001
- 3 with `2003
- 20 with `1920
- 20052009 with `2005
- 215 with `2015
- 71 with `1971
- 4965 with `1965

```
In [10]: df['build_year'].replace(0.0, 2000.0, inplace=True)
df['build_year'].replace(1.0, 2001.0, inplace=True)
df['build_year'].replace(3.0, 2003.0, inplace=True)
df['build_year'].replace(20.0, 1920.0, inplace=True)
df['build_year'].replace(20052009.0, 2005.0, inplace=True)
df['build_year'].replace(215.0, 2015.0, inplace=True)
df['build_year'].replace(71.0, 1971.0, inplace=True)
df['build_year'].replace(4965.0, 1965.0, inplace=True)
```

life_sq

```
In [11]: df['life_sq'].plot.box(figsize=(15,10), grid=True);
```



The boxplot shows a clear outlier with a value above 7000, which we can simply filter (we want to keep 'NaN'):

```
In [12]: mask = (df['life_sq'] < 1000) | (df['life_sq'].isnull())
df = df[mask]
```

Imputation with Mean

```
In [13]: df_mean = df.copy()
```

```
In [14]: mean = df_mean['life_sq'].mean(); mean
```

```
Out[14]: 34.09424170714493
```

```
In [15]: df_mean['life_sq'].fillna(mean, inplace=True)
```

```
In [16]: df_mean['life_sq'].isnull().sum()
```

```
Out[16]: 0
```

KNN-Imputation

Since we cannot use all features for KNN imputation for efficiency reasons, we want to select features that correlate with `life_sq`.

```
In [17]: df_knn = df.select_dtypes(include=[np.number])
```

```
In [18]: df_knn.dtypes.unique()
```

```
Out[18]: array([dtype('int64'), dtype('float64')], dtype=object)
```

```
In [19]: correlations = {}
         for i, feature in enumerate(df_knn.columns):
             correlations[i] = df_knn['life_sq'].corr(df_knn[feature])

In [20]: sorted(correlations.items(), key=lambda x: x[1], reverse=True)[:10]

Out[20]: [(2, 1.0),
          (7, 0.5513954768950734),
          (275, 0.4148448606438824),
          (1, 0.3951990153935231),
          (6, 0.2198999416737205),
          (125, 0.1508240306510387),
          (10, 0.14751393386330344),
          (261, 0.14549050222571336),
          (259, 0.14541745192012867),
          (260, 0.14533108257930286)]
```

We still need to check how many NaN values these features have. At best, these should only be a few...

```
In [21]: df_knn.iloc[:, [2, 7, 275, 1, 6, 125, 10, 261, 259, 260]].isnull().sum()

Out[21]: life_sq                6383
         num_room             9572
         price_doc              0
         full_sq              0
         build_year          13605
         workplaces_km         0
         area_m               0
         cafe_avg_price_5000    297
         cafe_sum_5000_min_price_avg 297
         cafe_sum_5000_max_price_avg 297
         dtype: int64
```

Apart from `build_year`, many values are available, so we can perform the imputation in this way.

```
In [22]: from sklearn.impute import KNNImputer

In [23]: imputer = KNNImputer(n_neighbors=2)

In [24]: imputed = imputer.fit_transform(df_knn.iloc[:, [2, 7, 275, 1, 6, 125, 10, 261, 259]
```

We create a new DataFrame with the filled values:

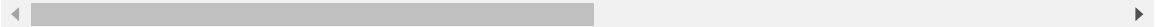
```
In [25]: df_imputed = pd.DataFrame(imputed, columns=df_knn.iloc[:, [2, 7, 275, 1, 6, 125]

Check whether all values have been replaced:

In [26]: df_imputed.head()
```

Out[26]:

	life_sq	num_room	price_doc	full_sq	build_year	workplaces_km	area_m
0	27.0	1.0	5850000.0	43.0	1986.0	0.884350	6.407578e+06
1	19.0	1.5	6000000.0	34.0	1972.5	0.686252	9.589337e+06
2	29.0	1.5	5700000.0	43.0	1963.5	1.510089	4.808270e+06
3	50.0	4.0	13100000.0	89.0	2014.5	0.622272	1.258354e+07
4	77.0	3.5	16331452.0	77.0	1947.5	0.892668	8.398461e+06



In [27]: `df_imputed.isnull().sum()`

Out[27]:

life_sq	0
num_room	0
price_doc	0
full_sq	0
build_year	0
workplaces_km	0
area_m	0
cafe_avg_price_5000	0
cafe_sum_5000_min_price_avg	0
cafe_sum_5000_max_price_avg	0

dtype: int64