

Introduction to Machine Learning – AICOSS

Machine learning – the *why* and the *what*

Why machine learning?



„The world’s most valuable resource is no longer oil, but data“

The Economist 2017, <https://www.economist.com/leaders/2017/05/06/the-worlds-most-valuable-resource-is-no-longer-oil-but-data>

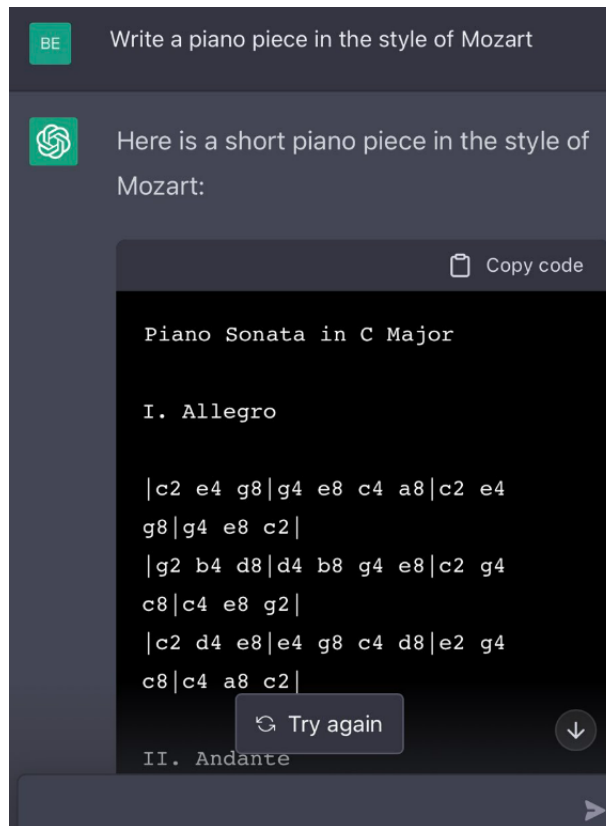
Why machine learning?




ChatGPT



Stable Diffusion



„christmas tree, forest, moon, squirrel, mystical fog,hr giger“



Rombach et al., *High-Resolution Image Synthesis with Latent Diffusion Models*, 2022


Why machine learning?

| | |
|-----------------------------|---|
| Natural Language Processing | Question answering; speech recognition; summarizing documents; classifying documents; finding entities in documents; searching documents for concepts |
| Computer Vision | Interpreting satellite and drone images (disaster detection); face recognition; automatic image description; recognizing traffic signs; autonomous driving |
| Medicine | Recognize anomalies in radiological images (CT, MRI and X-ray); tissue classification in microscopy images; feature identification in ultrasound |
| Biology | Protein folding; classification of proteins; applications in genomics (e.g. tumor-normal sequencing); cell classification; analysis of protein/protein interactions |
| Image generation | Colorizing B&W images; increasing image resolution; removing noise from images; transferring stylistics to images (e.g. Instagram filters) |
| Recommender Systems | Web search; Product recommendations |
| Games | Superhuman performance in chess, Go, many Atari games and real-time strategy games |
| Robotics | Dealing with objects that are difficult to locate (transparent, shiny, little texture) or that are generally difficult to grasp |
| Other applications | Finance and logistics forecasting; text-to-speech; many many more... |

Why machine learning?

Practical reasons


- Less programming; more efficient
- Scaling and customization of products
- Solving "non-programmable" tasks



„Philosophical“
reasons


What is machine learning?

Classical computer science approach:




What is machine learning?

Classical computer science approach:




What is machine learning?

Classical computer science approach:




Machine Learning:




What is machine learning?

Testing:



Training:



What is machine learning?

Formal definition: A computer program **A** is said to learn from experience **E** with respect to some class of tasks **T** and performance measure **P**, if its performance at tasks in **T**, as measured by **P**, improves with experience **E**. (*Tom M. Mitchell, 1997*)

Informal : Algorithms that *improve* in a *task* with increasing *experience*.

ML History

ML History

Samuel's Checkers Playing Program (1952)


- Program improves with the number of games (machine learning...)
- Evaluation of the board positions by a search tree
- Best move is determined using a minimax strategy



ML History

Perceptron 1957 (Frank Rosenblatt)


- Predecessor of today's neural networks:
Multilayer Perceptron (MLP) -> Artificial Neural Networks -> Deep Learning
- Triggered great excitement around AI ("AI Boom")



ML History

AI winter

- Minsky & Papert (1969) "ruin" AI
- Public and private investment in AI technology and research collapses for decades




ML History

Resurrection as machine learning

Differences:

- ML: bottom-up, AI: top-down
- ML: More practical application, based on statistics and optimization, not logic



ML History

TD-Gammon (Gerald Tesauro 1992)

- Gerry Tesauro (IBM) trains a neural network that plays against itself again and again (100k+). It beats several backgammon world champions.



ML History

IBM Deep Blue (1997)

- IBM's Deep Blue chess computer wins against Garry Kasparov.
- Predominantly "brute-force" approach with calculation of 126 million positions per second.
But the evaluation of a position is ML.



<https://www.ibm.com/blogs/think/2017/05/deep-blue/>

ML History

AI today: „Deep Learning Revolution“

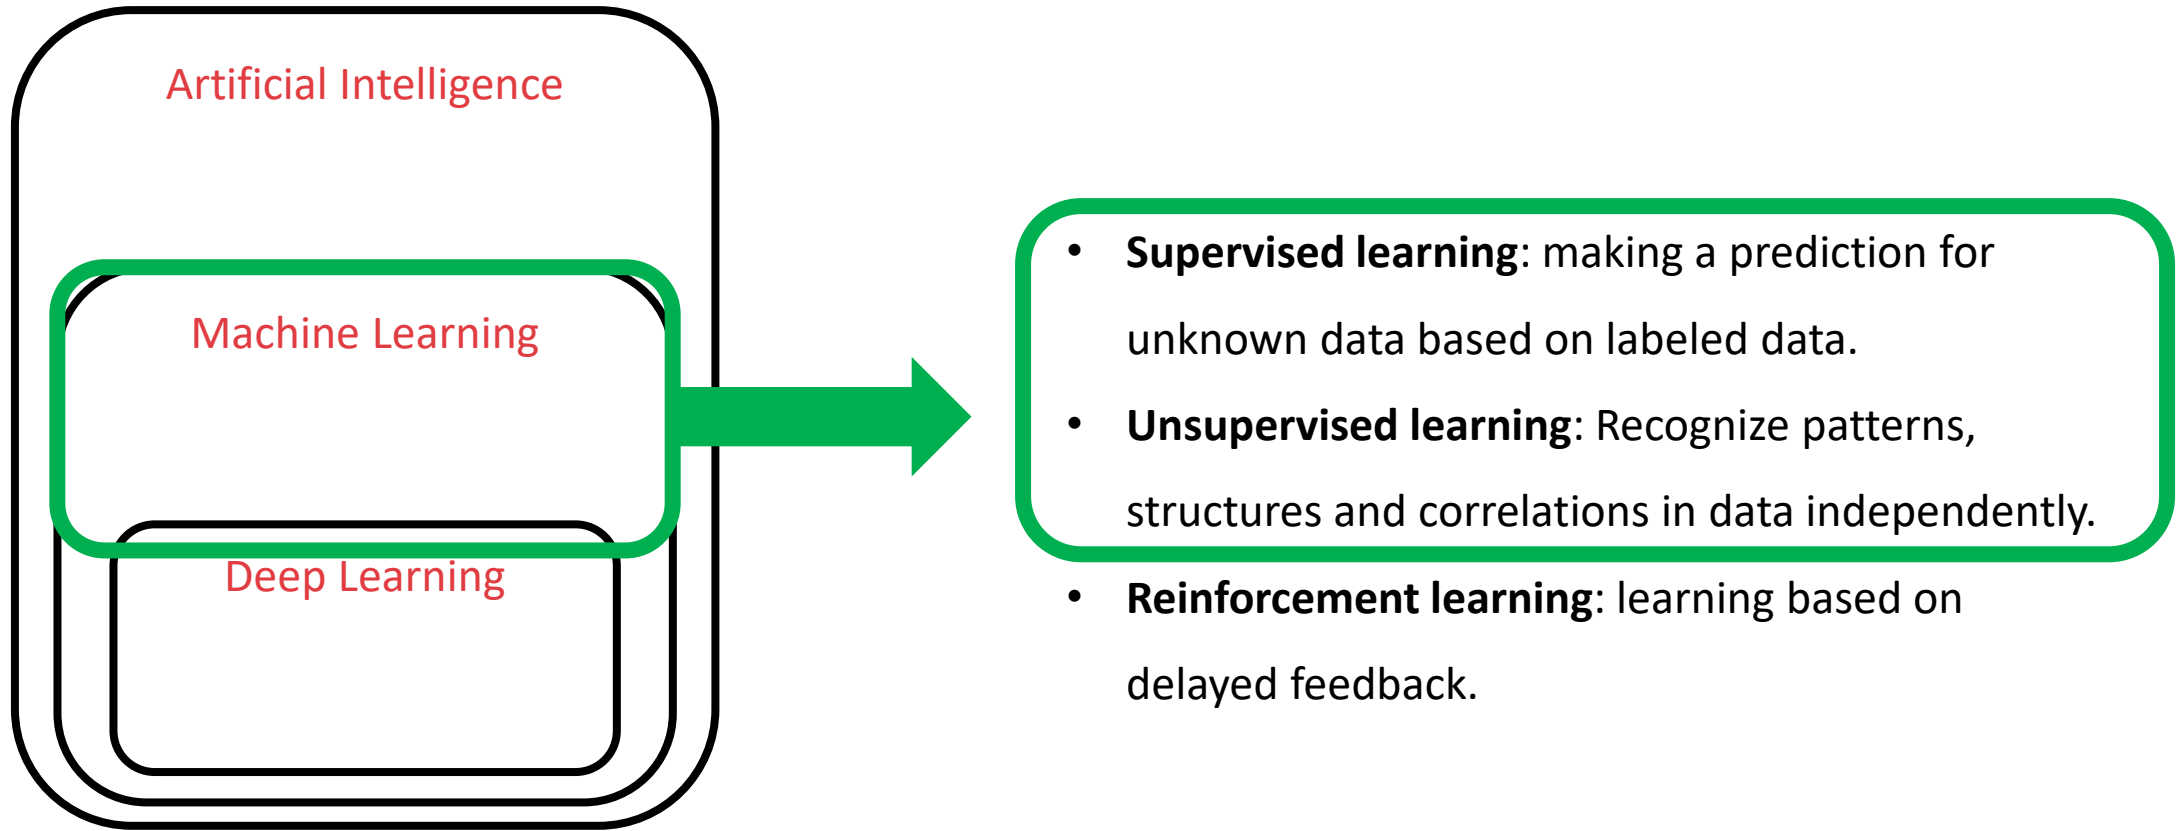
- AlexNET wins 2012 ImageNet competition
- AlphaGo beats Lee Sedol 2016 in Go (2.1×10^{170} positions, cf. chess 10^{40})
- Now: LLMs, self-driving cars, robotics, etc.



<https://www.theguardian.com/technology/2016/mar/15/alphago-what-does-google-advanced-software-go-next>


Scope and structure of this week

ML scope



The **application** is the central element of the course!


Structure of this weeks course



Contents

- Python for data science recap
- Unsupervised learning
- Supervised learning
- Evaluation and hyperparameter tuning
- Ensemble learning
- (Simple) neural networks


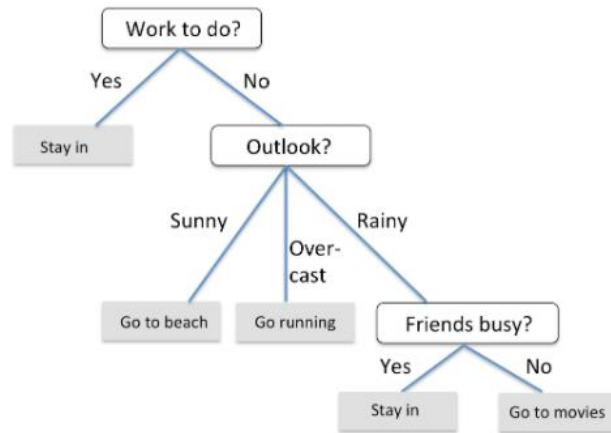
Structure of this weeks course



Contents

- Python for data science recap
- Unsupervised learning
- Supervised learning
- Evaluation and hyperparameter tuning
- Ensemble learning
- (Simple) neural networks

Structure of this weeks course





Contents

- Python for data science recap
- Unsupervised learning
- Supervised learning
 - Evaluation and hyperparameter tuning
 - Ensemble learning
 - (Simple) neural networks

Structure of this weeks course

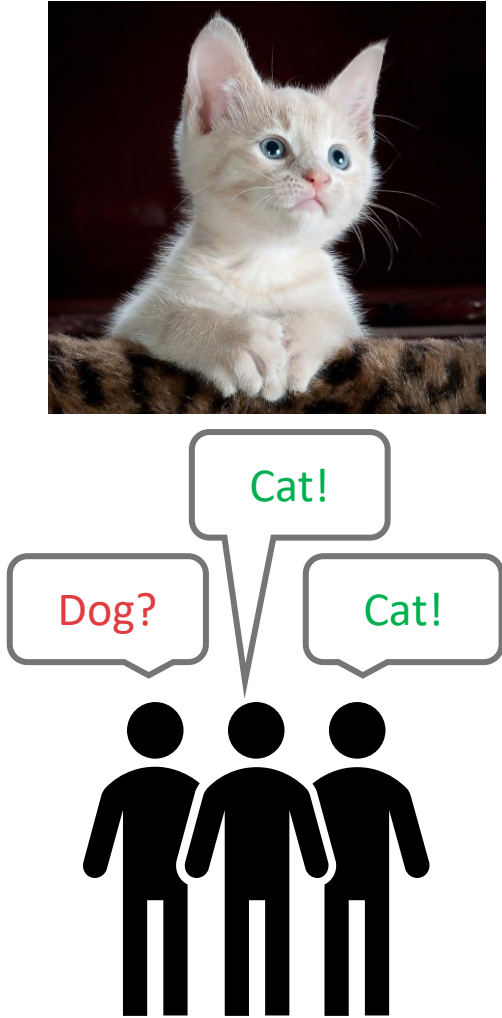
| | | Predicted class | |
|--------------|---|----------------------|----------------------|
| | | P | N |
| Actual class | P | True positives (TP) | False negatives (FN) |
| | N | False positives (FP) | True negatives (TN) |



Contents

- Python for data science recap
- Unsupervised learning
- Supervised learning
- Evaluation and hyperparameter tuning
- Ensemble learning
- (Simple) neural networks

Structure of this weeks course




Contents

- Python for data science recap
- Unsupervised learning
- Supervised learning
- Evaluation and hyperparameter tuning
- Ensemble learning
- (Simple) neural networks

Structure of this weeks course

Contents



- Python for data science recap
- Unsupervised learning
- Supervised learning
- Evaluation and hyperparameter tuning
- Ensemble learning
- (Simple) neural networks